

# hkvrbhiuu

January 23, 2025

```
[44]: import pandas as pd
from sklearn.model_selection import train_test_split
import xgboost as xgb
from sklearn.metrics import accuracy_score, classification_report, \
    confusion_matrix, mean_squared_error, mean_absolute_error, r2_score
import matplotlib.pyplot as plt
import seaborn as sns
import math
```

```
[45]: # Load Titanic dataset
data = pd.read_csv(r"C:\Users\91703\OneDrive\Desktop\TITANIC.csv")
```

```
[46]: data.head()
```

```
[46]:
```

	PassengerId	Survived	Pclass	\
0	1	0	3	
1	2	1	1	
2	3	1	3	
3	4	1	1	
4	5	0	3	

	Name	Sex	Age	SibSp	\
0	Braund, Mr. Owen Harris	male	22.0	1	
1	Cumings, Mrs. John Bradley (Florence Briggs Th...	female	38.0	1	
2	Heikkinen, Miss. Laina	female	26.0	0	
3	Futrelle, Mrs. Jacques Heath (Lily May Peel)	female	35.0	1	
4	Allen, Mr. William Henry	male	35.0	0	

	Parch	Ticket	Fare	Cabin	Embarked
0	0	A/5 21171	7.2500	NaN	S
1	0	PC 17599	71.2833	C85	C
2	0	STON/O2. 3101282	7.9250	NaN	S
3	0	113803	53.1000	C123	S
4	0	373450	8.0500	NaN	S

```
[47]: data.describe()
```

```
[47]:
```

	PassengerId	Survived	Pclass	Age	SibSp \
count	891.000000	891.000000	891.000000	714.000000	891.000000
mean	446.000000	0.383838	2.308642	29.699118	0.523008
std	257.353842	0.486592	0.836071	14.526497	1.102743
min	1.000000	0.000000	1.000000	0.420000	0.000000
25%	223.500000	0.000000	2.000000	20.125000	0.000000
50%	446.000000	0.000000	3.000000	28.000000	0.000000
75%	668.500000	1.000000	3.000000	38.000000	1.000000
max	891.000000	1.000000	3.000000	80.000000	8.000000

	Parch	Fare
count	891.000000	891.000000
mean	0.381594	32.204208
std	0.806057	49.693429
min	0.000000	0.000000
25%	0.000000	7.910400
50%	0.000000	14.454200
75%	0.000000	31.000000
max	6.000000	512.329200

```
[48]: data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 891 entries, 0 to 890
Data columns (total 12 columns):
#   Column          Non-Null Count  Dtype
---  -
0   PassengerId      891 non-null    int64
1   Survived         891 non-null    int64
2   Pclass           891 non-null    int64
3   Name             891 non-null    object
4   Sex              891 non-null    object
5   Age              714 non-null    float64
6   SibSp            891 non-null    int64
7   Parch            891 non-null    int64
8   Ticket           891 non-null    object
9   Fare             891 non-null    float64
10  Cabin            204 non-null    object
11  Embarked         889 non-null    object
dtypes: float64(2), int64(5), object(5)
memory usage: 83.7+ KB
```

```
[49]: # Select relevant columns and convert categorical variables
data = data[['Survived', 'Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare']]
data['Sex'] = data['Sex'].map({'male': 0, 'female': 1})
```

```
[50]: data.head()
```

```
[50]:
```

	Survived	Pclass	Sex	Age	SibSp	Parch	Fare
0	0	3	0	22.0	1	0	7.2500
1	1	1	1	38.0	1	0	71.2833
2	1	3	1	26.0	0	0	7.9250
3	1	1	1	35.0	1	0	53.1000
4	0	3	0	35.0	0	0	8.0500

```
[51]: # Handle missing Age values
data['Age'] = pd.to_numeric(data['Age'], errors='coerce')
data['Age'] = data['Age'].fillna(data['Age'].mean())
```

```
[52]: # Define target variable (y) and feature variables (X)
y = data['Survived']
X = data[['Pclass', 'Sex', 'Age', 'SibSp', 'Parch', 'Fare']]
```

```
[53]: # Split data into training and testing sets
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳random_state=42)
```

```
[54]: # Create and train XGBoost Classifier model
model = xgb.XGBClassifier(objective='binary:logistic', n_estimators=100,
↳random_state=42)
model.fit(X_train, y_train)
```

```
[54]: XGBClassifier(base_score=None, booster=None, callbacks=None,
      colsample_bylevel=None, colsample_bynode=None,
      colsample_bytree=None, device=None, early_stopping_rounds=None,
      enable_categorical=False, eval_metric=None, feature_types=None,
      gamma=None, grow_policy=None, importance_type=None,
      interaction_constraints=None, learning_rate=None, max_bin=None,
      max_cat_threshold=None, max_cat_to_onehot=None,
      max_delta_step=None, max_depth=None, max_leaves=None,
      min_child_weight=None, missing=nan, monotone_constraints=None,
      multi_strategy=None, n_estimators=100, n_jobs=None,
      num_parallel_tree=None, random_state=42, ...)
```

```
[55]: # Make predictions on test data
y_pred = model.predict(X_test)
print(y_pred)
```

```
[0 0 1 1 1 1 1 0 0 1 1 0 0 0 0 1 1 1 0 0 0 1 0 0 0 0 0 0 0 1 0 1 1 1 0 0 0
 1 1 0 0 0 0 0 0 0 0 0 0 0 1 1 1 0 1 0 1 0 1 1 0 0 1 1 0 0 1 0 1 0 1 1 1 1
 0 0 1 1 1 1 0 1 1 0 1 0 1 1 0 0 1 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 1 0 0 0 1
 0 1 1 0 0 0 0 1 0 0 1 1 1 0 1 1 0 0 0 1 0 0 1 1 0 1 1 0 0 0 1 0 0 0 1 0 0
 1 0 0 0 0 0 0 0 0 1 1 1 0 0 0 1 0 0 0 1 0 0 0 1 1 1 0 0 0 1 1]
```


```
[56]: # Evaluate model performance
accuracy = accuracy_score(y_test, y_pred)
print("Accuracy:", accuracy)
print("Classification Report:", classification_report(y_test, y_pred))
print("Confusion Matrix:", confusion_matrix(y_test, y_pred))
```

Accuracy: 0.8044692737430168

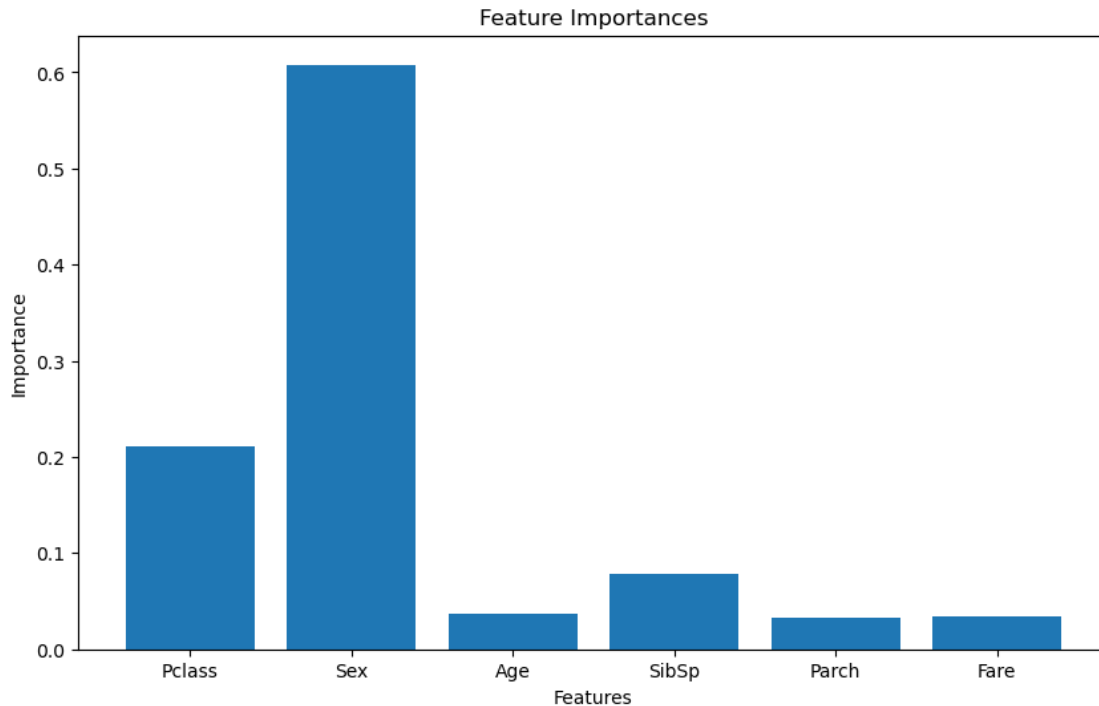
Classification Report:		precision	recall	f1-score	support
	0	0.83	0.84	0.83	105
	1	0.77	0.76	0.76	74
	accuracy			0.80	179
	macro avg	0.80	0.80	0.80	179
	weighted avg	0.80	0.80	0.80	179

Confusion Matrix: [[88 17]  
[18 56]]

```
[57]: # Plot Confusion Matrix
plt.figure(figsize=(8, 6))
sns.heatmap(confusion_matrix(y_test, y_pred), annot=True, cmap='Blues')
plt.xlabel("Predicted labels")
plt.ylabel("True labels")
plt.show()
```



```
[58]: #Plot Feature Importances  
feature_importances = model.feature_importances_  
plt.figure(figsize=(10, 6))  
plt.bar(X.columns, feature_importances)  
plt.xlabel("Features")  
plt.ylabel("Importance")  
plt.title("Feature Importances")  
plt.show()
```



```
[59]: # Define target variable (y) and feature variables (X)
```

```
y = data['Age']
X = data[['Pclass', 'Sex', 'SibSp', 'Parch', 'Fare']]
```

```
[60]: # Split data into training and testing sets
```

```
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
↳ random_state=42)
```

```
[61]: # Create and train XGBoost Regressor model
```

```
model = xgb.XGBRegressor(objective='reg:squarederror', n_estimators=100,
↳ random_state=42)
model.fit(X_train, y_train)
```

```
[61]: XGBRegressor(base_score=None, booster=None, callbacks=None,
    colsample_bylevel=None, colsample_bynode=None,
    colsample_bytrees=None, device=None, early_stopping_rounds=None,
    enable_categorical=False, eval_metric=None, feature_types=None,
    gamma=None, grow_policy=None, importance_type=None,
    interaction_constraints=None, learning_rate=None, max_bin=None,
    max_cat_threshold=None, max_cat_to_onehot=None,
    max_delta_step=None, max_depth=None, max_leaves=None,
    min_child_weight=None, missing=nan, monotone_constraints=None,
    multi_strategy=None, n_estimators=100, n_jobs=None,
    num_parallel_tree=None, random_state=42, ...)
```

```
[62]: # Make predictions on test data
y_pred = model.predict(X_test)
print(y_pred)
```

```
[34.456184  32.783123  31.337349  27.583454  12.894008  31.473656
 29.464012  19.475527  29.464012  28.124666  24.621515  29.840971
 29.431765  31.094555  31.998625  47.26982   35.157482  21.611671
 31.998625  32.23708   29.390882  50.091675  19.80008   19.897299
 25.972647   9.876298  30.707668  31.998625   9.876298  21.01747
 29.390882  29.464012  39.301514  23.321783  29.840971  26.733416
 42.585773  29.464012  25.319326  29.390882  27.743006  24.906473
 29.840971  33.80651   17.985836  48.399864  28.493532  19.897299
 24.202341  24.283619   0.80724084 43.92668   48.18969   35.89958
 33.80651   29.847437  32.783123  31.010103  30.077312  28.611385
 30.548288  40.837154  28.281397  46.406418  33.80651   35.398415
 18.429457  28.295893  14.822876  51.58079   7.578495  21.083878
 44.3008    49.961662  28.295893  30.643974  29.464012  27.291847
 31.505043  23.677633   5.9831247 35.398415  49.754482  28.022041
 39.301514  28.525219  22.62831   28.47828   34.623486  32.047707
 24.633507   5.954818  39.014965  28.621103  29.390882  28.295893
 35.83145   29.840971  31.505043  29.390882  33.403744  31.512476
 30.90827   24.202341  29.840971  25.321844  51.772957  33.717537
 26.043966  32.23393   35.828175  32.047707  39.704784  39.02451
 31.463264  20.308886  42.5742    30.98962   26.874895  48.461018
 0.69820094 34.785942  50.203407  17.815742  32.047707  50.091675
 47.82472   43.08586   21.743963  29.840971  28.611385  34.456184
 29.83419   29.948875  39.568726  22.571812  24.792702  30.548288
 29.179548  30.314743  33.66125   28.281397  31.998625  36.125763
 34.11547   22.282433  31.998625  32.783123  49.59976   29.840971
 29.9259    31.956175  24.571924  28.51809   29.840971  24.202341
 29.59447   22.045774  28.281397  29.808167  28.21134   50.091675
 31.998625  25.184313  33.66125   50.157597  32.783123  51.962532
 29.390882  24.202341  50.157597  18.922506  39.02451   29.808167
 31.512476  31.094555  40.300625  35.398415  26.42828   ]
```

```
[63]: # Evaluate model performance
mse = mean_squared_error(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
rmse = math.sqrt(mse)
r2 = r2_score(y_test, y_pred)
print("Mean Squared Error:", mse)
print("Mean Absolute Error:", mae)
print("Root Mean Squared Error:", rmse)
print("R-squared:", r2)
```

```
Mean Squared Error: 154.5745228332585
Mean Absolute Error: 9.080691027556679
Root Mean Squared Error: 12.432800281242296
```

R-squared: 0.08703754490610083

```
[64]: # Plot predicted vs actual values
plt.figure(figsize=(10, 6))
plt.scatter(y_test, y_pred)
plt.xlabel("Actual Age")
plt.ylabel("Predicted Age")
plt.show()
```

